

Capstone Project Kickoff

Term Project (Units 1–3 Integrated, emphasis on Unit 3)

Weight: 30% of course grade

Milestones: Proposal Tue, November 25;
Milestone demo Week 15 Thu;
Final deliverables Week 16 Thu

Datasets: Team-chosen **Kaggle** dataset (primary **batch** source)
Plus one streaming public API (secondary live source)

Public APIs Free and Open Public APIs (No Auth Needed)

<https://mixedanalytics.com/blog/list-actually-free-open-no-auth-needed-apis/>

Tools: Colab, GitHub, GCS, BigQuery, **Cloud Functions 2nd gen**, **Pub/Sub**, **Dataflow template**, **BQML**, Looker Studio, Gemini prompts

Minimum Requirements

1. Batch Ingest (from Kaggle)

- Store raw data in **GCS**, load curated tables into **BigQuery** (document schema & partitioning).
- Provide **one** data quality check and **one** transformation logic explanation.

2. Streaming Ingest (from API)

- **Cloud Function (2nd gen)** producer that calls a public API periodically (e.g., weather, crypto, transit, air quality, finance; no auth secrets in code).
- Publish payloads to **Pub/Sub** as normalized JSON.
- **Pipeline (API → Pub/Sub → BigQuery)** template writing to a streaming fact table.
- Validate with queries showing near real-time rows (timestamps).

Minimum Requirements

3. Analytics & Modeling

- **At least one** BQML model (regression or classification) that uses **both** batch and streaming features (e.g., join last N streaming metrics to batch entities, or score batch data with a streaming covariate).
- **ML.EVALUATE** and **ML.EXPLAIN_PREDICT** examples; threshold discussion if classification.

4. Executive Dashboard

- Looker Studio OR Plotly dashboard with **3–5 KPIs** and **at least one** time-series component fed by the streaming table.

Minimum Requirements

5. Prompt Engineering & DIVE

- Include **prompt design** for analysis and **DIVE** journal entries showing how prompts evolved, where they failed, and how you validated.

6. Architecture + Ops

- **Diagram** (GCS → BQ; Function → Pub/Sub → BQ; BQML; Looker/Plotly).
- Reproducibility instructions to spin up/down the pipeline.

Milestones

- **Nov. 25 (Tue): Blueprint** (as in Lab 9) – business problem, sources, architecture, ML plan, KPIs, risks.
- **Dec. 11 (Thu): Live pipeline demo** (Cloud Function invoke + Pub/Sub → BigQuery row appears; show dashboard refresh).
- **Dec. 13 (Sun): Final presentation + code freeze.**

Deliverables

- **Individual (60 points total):**
 - **Notebook:** Final_<Name>_analysis.ipynb
 - Prompt logs, DIVE entries for one substantive question, at least **one** interactive Plotly figure, and a link to the dashboard section they influenced.
 - **Contribution MD (1 page):** Final_<Name>_contrib.md (exact tasks, PR links, lessons learned).
- **Team (40 points total):**
 - **Architecture README:** how to deploy/teardown; service enablement; IaC checklist (gcloud commands OK).
 - **Live Demo (8–10 slides + 5-min demo):** problem, architecture, pipeline proof, model results, dashboard, business impact.
 - **Governance Note (1–2 pages):** assumptions, data ethics, privacy/security notes, failure playbook.

Submission: D2L ZIP + repo URL + Looker link. Include **two screenshots:** BQ table screenshot with latest rows.

Suggested Repo Layout

```
TermProject_TeamX/
├── notebooks/
│   ├── Final_Alice_analysis.ipynb
│   ├── Final_Bob_analysis.ipynb
│   └── ...
├── pipeline/
│   ├── function/      (main.py, requirements.txt)
│   └── infra/         (gcloud commands, enable APIs, SA roles)
├── bq/
│   └── sql/           (queries for eval/predict/join)
├── dashboards/
│   └── kpis.md        (definitions + Looker link)
└── docs/
    ├── blueprint.pdf
    ├── governance.pdf
    └── ops_runbook.md
 README.md
```

Rubric (100 + up to 10 EC)

Individual (60)

- DIVE rigor (clear pivot from first to validated insight) – **20**
- Analytical depth + correct use of prompts & Plotly – **20**
- Clear, reproducible notebook & contribution mapping – **20**

Team (40)

- Pipeline completeness (batch + streaming) & validation – **15**
- BQML relevance + evaluation/explainability – **10**
- Dashboard clarity & KPI correctness – **10**
- Cost/ops considerations – **5**

Extra Credit (up to +10)

- Machine learning model quality

Unit 3 - Lab 2

OpenSky to BigQuery Table

Building Data Pipelines with BigQuery ML



2. Google Cloud Storage

Data Flow

Upload data to GCS bucket for staging



3. BigQuery Table

Load data into BigQuery table for analysis



4. BQML Regression

ML Model

Predict flight velocity using linear regression





4. BQML Regression ML Model

Predict flight velocity using linear regression



5. BQML Classification ML Model

Classify if flight is on ground



Critical Learning: NULL Handling

The classification model initially failed because the WHERE clause filtered out all 383 on_ground=TRUE records (which had NULL altitude/velocity values).

Solution: Use `COALESCE(altitude, 0)` to handle NULLs and preserve label diversity.

Cell 1: Python Packages and Authentication

Purpose

This cell installs required Python packages and authenticates the user to Google Cloud Platform (GCP).

Packages Installed

- **google-cloud-storage**: For interacting with Google Cloud Storage
- **google-cloud-bigquery**: For interacting with BigQuery
- **requests**: Popular HTTP library for API calls

Note: Authenticates user to enable access to GCP services

Cell 2: Project Configuration

Purpose

Configures essential project-specific variables for Google Cloud operations and sets the gcloud project.

Variables Defined

- **PROJECT_ID:** GCP project identifier
- **GCP_REGION:** Cloud region for services
- **GCS_BUCKET_NAME:** Cloud Storage bucket name
- **GCS_FOLDER_PATH:** Data storage folder path
- **BQ_DATASET:** BigQuery dataset name
- **BQ_TABLE:** BigQuery table for flight data
- **FLIGHT_RECORD_LIMIT:** API record limit

Cell 3: OpenSky API Class

Purpose

Defines the OpenSkyApi class and helper functions for data parsing and formatting.

Key Components

- OpenSky API client class implementation
- Data parsing functions for flight information
- Data formatting utilities for BigQuery compatibility

Cell 4: Initialize GCP Clients & Schema

Purpose

Initializes GCP clients, defines BigQuery schema, and implements the complete data pipeline logic.

Key Activities

- Initialize Google Cloud Storage and BigQuery clients
- Define BigQuery table schema for flight data
- Implement pipeline logic (API → GCS → BigQuery)

Cell 5: Execute Data Pipeline

Purpose

Executes the complete data pipeline from API to BigQuery.

Pipeline Flow

Fetch Data → Upload to GCS → Load into BigQuery

Cell 6: Full Pipeline Orchestration

Purpose

Orchestrates the complete end-to-end data pipeline from API ingestion through to BigQuery storage.

- Coordinates API data fetching
- Manages GCS upload operations
- Handles BigQuery table loading

Cell 7: BQML Regression Model

Purpose

Creates a BigQuery ML regression model to predict flight velocity.

Model Details

- **Type:** Linear Regression
- **Target:** Flight velocity
- **Features:** Altitude, latitude, longitude

Cell 8: Label Diversity Analysis

Purpose

Analyzes the on_ground label diversity in BigQuery to troubleshoot classification model issues.

Investigation Focus

- Examines distribution of on_ground values
- Identifies NULL handling issues
- Determines why initial model training failed

Cell 9: BQML Classification Model

Purpose

Creates a BigQuery ML logistic regression model to classify whether a flight is on the ground.

Model Specifications

- **Type:** Logistic Regression
- **Label:** on_ground (boolean)
- **Features:** altitude, velocity
- **Fix:** Uses COALESCE for NULL handling

Lab Summary: Key Learnings

Primary Issue & Resolution

Classification model failed due to WHERE clause filtering out TRUE labels. Fixed with COALESCE NULL handling.

Key Findings

- 383 on_ground=TRUE records had NULL values
- Successfully created regression and classification models

Best Practices

- Verify data after applying filters
- Use COALESCE for NULL handling